


ID: 18521272

Name: Le Ngoc Thai Phuong

I. Classification

```
In [1]: 1 from sklearn import datasets
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6
7 digits = datasets.load_digits()
8
9 # Display digit 1010
10 plt.imshow(digits.images[1010], cmap=plt.cm.gray_r, interpolation='nearest')
11 plt.show()
```



```
In [2]: 1 from sklearn.neighbors import KNeighborsClassifier
2 from sklearn.model_selection import train_test_split
3
4 # Create feature and target arrays
5 X = digits.data
6 y = digits.target
7
8 # Split into training and test set
9 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
10 random_state=42, stratify=y)
11
12 # Create a k-NN classifier with 7 neighbors: knn
13 knn = KNeighborsClassifier(n_neighbors=7)
14
15 # Fit the classifier to the training data
16 knn.fit(X_train, y_train)
17
18 # Print the accuracy
19 print(knn.score(X_test, y_test))
```

0.9833333333333333

C:\Users\ACER\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```

In [3]: 1 neighbors = np.arange(1, 9)
        2 train_accuracy = np.empty(len(neighbors))
        3 test_accuracy = np.empty(len(neighbors))
        4
        5 # Loop over different values of k
        6 for i, k in enumerate(neighbors):
        7     # Setup a k-NN Classifier with k neighbors: knn
        8     knn = KNeighborsClassifier(n_neighbors=k)
        9
        10    # Fit the classifier to the training data
        11    knn.fit(X_train, y_train)
        12
        13    # Compute accuracy on the training set
        14    train_accuracy[i] = knn.score(X_train, y_train)
        15
        16    # Compute accuracy on the testing set
        17    test_accuracy[i] = knn.score(X_test, y_test)
        18
        19 # Generate plot
        20 plt.title('k-NN: Varying Number of Neighbors')
        21 plt.plot(neighbors, test_accuracy, label = 'Testing Accuracy')
        22 plt.plot(neighbors, train_accuracy, label = 'Training Accuracy')
        23 plt.legend()
        24 plt.xlabel('Number of Neighbors')
        25 plt.ylabel('Accuracy')
        26 plt.show()

```

C:\Users\ACER\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\ACER\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\ACER\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value

```

In [4]: 1 from __future__ import print_function
        2 import torch
        3 import torch.nn as nn
        4 import torch.nn.functional as F
        5 from torch.autograd import Variable

```

In []:

1

In [5]:

```
1 from torchvision import datasets, transforms
2 mnist = datasets.MNIST(root='', train=True, download=True)
3
4 print("Number of training example: ", mnist.train_data.shape)
5 print("Image information ", mnist[0])
```

Number of training example: torch.Size([60000, 28, 28])


Image information (<PIL.Image.Image image mode=L size=28x28 at 0x1D1426948B0>, 5)

C:\Users\ACER\anaconda3\lib\site-packages\torchvision\datasets\mnist.py:75: UserWarning: train_data has been renamed data
warnings.warn("train_data has been renamed data")

In [6]:

```
1 import matplotlib.pyplot as plt
2 %matplotlib inline
3 plt.imshow(mnist[0][0])
```

Out[6]: <matplotlib.image.AxesImage at 0x1d14a23f730>



```
In [7]: 1 class Net(nn.Module):
2         def __init__(self):
3             super(Net, self).__init__()
4
5             self.fully = nn.Sequential(
6                 nn.Linear(28*28, 10)
7             )
8
9         def forward(self, x):
10            x = x.view([-1, 28*28])
11            x = self.fully(x)
12            x = F.log_softmax(x, dim=1)
13            return x
```

```
In [8]: 1 train_loader = torch.utils.data.DataLoader(datasets.MNIST(root=".", train=
2 test_loader = torch.utils.data.DataLoader(datasets.MNIST(root=".", train=
3
```

```
In [9]: 1 def train():
2         learning_rate = 1e-3
3         num_epochs = 3
4
5         net = Net()
6         optimizer = torch.optim.Adam(net.parameters(), lr=learning_rate)
7
8         for epoch in range(num_epochs):
9             for batch_idx, (data, target) in enumerate(train_loader):
10                 output = net(data)
11
12                 loss = F.nll_loss(output, target)
13                 optimizer.zero_grad()
14                 loss.backward()
15                 optimizer.step()
16
17                 if batch_idx % 100 == 0:
18                     print('Epoch = %f. Batch = %s. Loss = %s' % (epoch, batch_
19
20         return net
```

In [10]: 1 net = train()

```
Epoch = 0.000000. Batch = 0. Loss = 2.339271306991577
Epoch = 0.000000. Batch = 100. Loss = 0.80986487865448
Epoch = 0.000000. Batch = 200. Loss = 0.6029140949249268
Epoch = 0.000000. Batch = 300. Loss = 0.5262631177902222
Epoch = 0.000000. Batch = 400. Loss = 0.5180872082710266
Epoch = 0.000000. Batch = 500. Loss = 0.49103665351867676
Epoch = 0.000000. Batch = 600. Loss = 0.2785436511039734
Epoch = 0.000000. Batch = 700. Loss = 0.5144723653793335
Epoch = 0.000000. Batch = 800. Loss = 0.2813175618648529
Epoch = 0.000000. Batch = 900. Loss = 0.29325219988822937
Epoch = 1.000000. Batch = 0. Loss = 0.2884300947189331
Epoch = 1.000000. Batch = 100. Loss = 0.2784908413887024
Epoch = 1.000000. Batch = 200. Loss = 0.45401087403297424
Epoch = 1.000000. Batch = 300. Loss = 0.24643376469612122
Epoch = 1.000000. Batch = 400. Loss = 0.35517826676368713
Epoch = 1.000000. Batch = 500. Loss = 0.2768622040748596
Epoch = 1.000000. Batch = 600. Loss = 0.3907529413700104
Epoch = 1.000000. Batch = 700. Loss = 0.3144478499889374
Epoch = 1.000000. Batch = 800. Loss = 0.22636058926582336
Epoch = 1.000000. Batch = 900. Loss = 0.25695928931236267
Epoch = 2.000000. Batch = 0. Loss = 0.12280141562223434
Epoch = 2.000000. Batch = 100. Loss = 0.5356943607330322
Epoch = 2.000000. Batch = 200. Loss = 0.2567916214466095
Epoch = 2.000000. Batch = 300. Loss = 0.3874041736125946
Epoch = 2.000000. Batch = 400. Loss = 0.33864039182662964
Epoch = 2.000000. Batch = 500. Loss = 0.405193954706192
Epoch = 2.000000. Batch = 600. Loss = 0.1978750377893448
Epoch = 2.000000. Batch = 700. Loss = 0.2189231514930725
Epoch = 2.000000. Batch = 800. Loss = 0.5021193027496338
Epoch = 2.000000. Batch = 900. Loss = 0.3062554895877838
```

In [11]:

```
1 net.eval()
2 test_loss = 0
3 correct = 0
4 total = 0
5
6 for data, target in test_loader:
7     total += len(target)
8     output = net(data)
9     pred = output.max(1, keepdim=True)[1]
10    correct += target.eq(pred.view_as(target)).sum()
11
12 print("Correct out of %s" % total, correct.item())
13 print("Percentage accuracy", correct.item()*100/10000.)
```


Correct out of 10000 9228
Percentage accuracy 92.28

II. Linear Regression

```
In [12]: 1 import numpy as np
          2 import pandas as pd
          3 import matplotlib.pyplot as plt
          4 import seaborn as sns
          5
          6 df = pd.read_csv('gapminder.csv')
```

```
In [13]: 1 sns.heatmap(df.corr(), square=True, cmap='RdYlGn')
```


Out[13]: <AxesSubplot:>



```
In [14]: 1 from sklearn.linear_model import LinearRegression
2
3 # Create the regressor: reg
4 reg = LinearRegression()
5
6 X_fertility = df['fertility'].values.reshape(-1, 1)
7 y = df['life'].values.reshape(-1, 1)
8
9 X_train, X_test, y_train, y_test = train_test_split(X_fertility, y, test_
10
11 # Create the prediction space
12 prediction_space = np.linspace(min(X_fertility), max(X_fertility)).reshape
13
14 # Fit the model to the data
15 reg.fit(X_train, y_train)
16
17 # compute predictions over the prediction space: y_pred
18 y_pred = reg.predict(prediction_space)
19
20 # Print  $R^2$ 
21 print(reg.score(X_fertility, y))
22
23 # Plot regression line on scatter plot
24 sns.scatterplot(x='fertility', y='life', data=df)
25 plt.plot(prediction_space, y_pred, color='black', linewidth=3)
```

0.6162438752151919

Out[14]: [<matplotlib.lines.Line2D at 0x1d14a2af730>]




```
In [15]: 1 features = pd.read_csv('gapminder.csv')
2 df = pd.read_csv('gapminder.csv')
3 del features['life']
4 del features['Region']
5
6 y_life = df['life'].values.reshape(-1,1)
7 x_train, x_test, y_train, y_test = train_test_split(features, y_life, tes
8
9 reg_all = LinearRegression()
10 reg_all.fit(x_train, y_train)
11 print(reg_all.score(features, y_life))
```

0.8914651485793137

In []: 1

In []: 1


In []: 1

In []: 1

Linear Regression using PyTorch

```
In [16]: 1 import matplotlib.pyplot as plt
2 %matplotlib inline
3 import numpy as np
```

```
In [17]: 1 N = 10 # number of data points
2 m = .9
3 c = 1
4 x = np.linspace(0,2*np.pi,N)
5 y = m*x + c + np.random.normal(0,.3,x.shape)
6 plt.figure()
7 plt.plot(x,y,'o')
8 plt.xlabel('x')
9 plt.ylabel('y')
10 plt.title('2D data (#data = %d)' % N)
11 plt.show()
```



```
In [18]: 1 import torch
```

Dataset

```
In [19]: 1 from torch.utils.data import Dataset
2 class MyDataset(Dataset):
3     def __init__(self, x, y):
4         self.x = x
5         self.y = y
6
7     def __len__(self):
8         return len(self.x)
9
10    def __getitem__(self, idx):
11        sample = {
12            'feature': torch.tensor([1, self.x[idx]]),
13            'label': torch.tensor([self.y[idx]])}
14        return sample
```

```
In [20]: 1 dataset = MyDataset(x, y)
2 for i in range(len(dataset)):
3     sample = dataset[i]
4     print(i, sample['feature'], sample['label'])
```

```
0 tensor([1., 0.], dtype=torch.float64) tensor([0.9039], dtype=torch.float64)
1 tensor([1.0000, 0.6981], dtype=torch.float64) tensor([1.2929], dtype=torch.float64)
2 tensor([1.0000, 1.3963], dtype=torch.float64) tensor([1.8805], dtype=torch.float64)
3 tensor([1.0000, 2.0944], dtype=torch.float64) tensor([3.0595], dtype=torch.float64)
4 tensor([1.0000, 2.7925], dtype=torch.float64) tensor([3.3478], dtype=torch.float64)
5 tensor([1.0000, 3.4907], dtype=torch.float64) tensor([4.7624], dtype=torch.float64)
6 tensor([1.0000, 4.1888], dtype=torch.float64) tensor([4.8760], dtype=torch.float64)
7 tensor([1.0000, 4.8869], dtype=torch.float64) tensor([5.6290], dtype=torch.float64)
8 tensor([1.0000, 5.5851], dtype=torch.float64) tensor([6.0992], dtype=torch.float64)
9 tensor([1.0000, 6.2832], dtype=torch.float64) tensor([6.5297], dtype=torch.float64)
```

Dataloader

```
In [29]: 1 from torch.utils.data import DataLoader
2
3 dataset = MyDataset(x, y)
4 batch_size = 4
5 shuffle = True
6 num_workers = 4
7 dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=shuffle, num_workers=num_workers)
```

```
In [ ]: 1 import pprint as pp
2 for i_batch, samples in enumerate(dataloader):
3     print('\nbatch# = %s' % i_batch)
4     print('samples: ')
5     pp.pprint(samples)
```

Model

```
In [23]: 1 import torch.nn as nn
2 import torch.nn.functional as F
3 class MyModel(nn.Module):
4     def __init__(self, input_dim, output_dim):
5         super(MyModel, self).__init__()
6         self.linear = nn.Linear(input_dim, output_dim)
7
8     def forward(self, x):
9         out = self.linear(x)
10        return out
```

Setting a model for our problem

```
In [24]: 1 input_dim = 2
2 output_dim = 1
3
4 model = MyModel(input_dim, output_dim)
```

Cost function

```
In [25]: 1 cost = nn.MSELoss()
```

Minimizing the cost function


```
In [ ]: 1 num_epochs = 10
        2 l_rate = 0.01
        3 optimiser = torch.optim.SGD(model.parameters(), lr = l_rate)
        4
        5 dataset = MyDataset(x, y)
        6 batch_size = 4
        7 shuffle = True
        8 num_workers = 4
        9 training_sample_generator = DataLoader(dataset, batch_size=batch_size, sh
10
11 for epoch in range(num_epochs):
12     print('Epoch = %s' % epoch)
13     for batch_i, samples in enumerate(training_sample_generator):
14         predictions = model(samples['feature'])
15         error = cost(predictions, samples['label'])
16         print('\tBatch = %s, Error = %s' % (batch_i, error.item()))
17         optimiser.zero_grad()
18         error.backward()
19         optimiser.step()
```

Lets see how well the model has learnt the data

```
In [27]: 1 x_for_plotting = np.linspace(0, 2*np.pi, 1000)
        2 design_matrix = torch.tensor(np.vstack([np.ones(x_for_plotting.shape), x_
        3 print('Design matrix shape:', design_matrix.shape)
        4
        5 y_for_plotting = model.forward(design_matrix)
        6 print('y_for_plotting shape:', y_for_plotting.shape)
```

```
Design matrix shape: torch.Size([1000, 2])
y_for_plotting shape: torch.Size([1000, 1])
```

```
In [28]: 1 plt.figure()
2 plt.plot(x,y,'o')
3 plt.plot(x_for_plotting, y_for_plotting.data.numpy(), 'r-')
4 plt.xlabel('x')
5 plt.ylabel('y')
6 plt.title('2D data (#data = %d)' % N)
7 plt.show()
```



III. Recommendation Systems

```
In [ ]: 1
```

```
In [ ]: 1
```